

# Separating the EoR Signal with a Convolutional Denoising Autoencoder: A Deep-learning-based Method

Weitian Li,<sup>1</sup>★ Haiguang Xu,<sup>1,2</sup>★ Zhixian Ma,<sup>3</sup> Ruimin Zhu,<sup>4</sup> Dan Hu,<sup>1</sup> Zhenghao Zhu,<sup>1</sup> Junhua Gu,<sup>5</sup> Chenxi Shan,<sup>1</sup> Jie Zhu<sup>3</sup> and Xiang-Ping Wu<sup>5</sup>

<sup>1</sup>*School of Physics and Astronomy, Shanghai Jiao Tong University, 800 Dongchuan Road, Shanghai 200240, China*

<sup>2</sup>*Tsung-Dao Lee Institute / IFSA Collaborative Innovation Center, Shanghai Jiao Tong University, 800 Dongchuan Road, Shanghai 200240, China*

<sup>3</sup>*Department of Electronic Engineering, Shanghai Jiao Tong University, 800 Dongchuan Road, Shanghai 200240, China*

<sup>4</sup>*Department of Statistics, Northwestern University, 2006 Sheridan Road, Evanston, IL 60208, US*

<sup>5</sup>*National Astronomical Observatories, Chinese Academy of Sciences, 20A Datun Road, Beijing 100012, China*

Accepted XXX. Received YYY; in original form ZZZ

## ABSTRACT

When applying the foreground removal methods to uncover the faint cosmological signal from the epoch of reionization (EoR), the foreground spectra are assumed to be smooth. However, this assumption can be seriously violated in practice since the unresolved or mis-subtracted foreground sources, which are further complicated by the frequency-dependent beam effects of interferometers, will generate significant fluctuations along the frequency dimension. To address this issue, we propose a novel deep-learning-based method that uses a 9-layer convolutional denoising autoencoder (CDAE) to separate the EoR signal. After being trained on the SKA images simulated with realistic beam effects, the CDAE achieves excellent performance as the mean correlation coefficient ( $\bar{\rho}$ ) between the reconstructed and input EoR signals reaches  $0.929 \pm 0.045$ . In comparison, the two representative traditional methods, namely the polynomial fitting method and the continuous wavelet transform method, both have outstanding difficulties in uncovering the EoR signal, yielding only  $\bar{\rho}_{\text{poly}} = 0.296 \pm 0.121$  and  $\bar{\rho}_{\text{cwt}} = 0.198 \pm 0.160$ , respectively. We conclude that, by hierarchically learning sophisticated features through multiple convolutional layers, the CDAE is a powerful tool that can be used to overcome the complicated frequency-dependent beam effects and accurately separate the EoR signal, which exhibits the great potential of deep-learning-based methods in future EoR experiments.

**Key words:** methods: data analysis – techniques: interferometric – dark ages, reionization, first stars – radio continuum: general

## 1 INTRODUCTION

The 21 cm line emission of neutral hydrogen from the epoch of reionization (EoR) is regarded as a decisive probe to directly explore this stage (see Furlanetto 2016 for a review). To detect the 21 cm signal, which is believed to have been redshifted to the frequencies below 200 MHz, low-frequency radio interferometers such as the SKA (Koopmans et al. 2015) and its pathfinders have been built or under construction. The observational challenges, however, are immense due to complicated instrumental effects, ionospheric distortions, radio frequency interference, and the strong foreground contamination that overwhelms the EoR signal by about 4–5 orders of magnitude (see Morales & Wyithe 2010 for a review). Fortunately, in the frequency dimension the foreground contamination is expected to be intrinsically smooth, while the EoR signal fluctuates

rapidly on  $\lesssim$  MHz scales. This difference is the key characteristic exploited by many foreground removal methods in order to uncover the faint EoR signal (e.g., Wang et al. 2006; Jelić et al. 2008; Liu et al. 2009b; Harker et al. 2009; Chapman et al. 2012, 2013; Wang et al. 2013; Gu et al. 2013; Bonaldi & Brown 2015; Mertens et al. 2018).

However, the smoothness of the foreground spectra can be destroyed by the frequency-dependent beam effects, i.e., the variation of the point spread function (PSF) with frequencies that cannot be perfectly calibrated (Liu et al. 2009a). Because of the incomplete  $uv$  coverage, the PSF has a complicated profile consisting of a narrow peaky main lobe and a multitude of jagged side lobes with relative amplitudes of about 0.1 per cent that extend beyond the field of view (e.g., Liu et al. 2009a, their figures 1 and 3). A source that is unresolved or mis-subtracted (e.g., due to the limited field of view) during the CLEAN process leaves catastrophic residuals, the locations of which vary with the frequency since the angular position

★ E-mail: liweitianux@sjtu.edu.cn (WL); hgxu@sjtu.edu.cn (HX)

of a PSF side lobe is inversely proportional to the frequency. These effects lead to complicated residuals fluctuating along the frequency dimension, which cannot be correctly separated from the EoR signal by the traditional foreground removal methods that rely on the smoothness of the foreground spectra.

Given the complicated profiles and frequency-dependent variations of the PSF, it would be difficult to craft a practicable model for most, if not all, existing foreground removal methods to overcome the beam effects, even at the cost of extensive computation burden (e.g., [Lochner et al. 2015](#)). Therefore deep-learning-based methods, which can distil knowledge from the data to automatically refine the model, seem more feasible and appealing (e.g., [Herbel et al. 2018](#); [Vafaei Sadr et al. 2018](#)). In recent years, the deep learning algorithms have seen prosperous developments and have brought breakthroughs into many fields (see [LeCun et al. 2015](#) for a recent review). Among various deep learning algorithms, the convolutional denoising autoencoder (CDAE) and its variants are flexible and powerful in learning subtle and complicated features from the data and have been successfully applied to weak gravitational wave signal denoising (e.g., [Shen et al. 2017](#)), monaural audio source separation (e.g., [Grais & Plumbley 2017](#)), and so on. These applications have demonstrated the outstanding abilities of the CDAE in extracting weak signals from highly temporal-variable data, thus it is worth trying to apply the CDAE to separate the EoR signal. Although the signal-to-noise ratio in the EoR detection is much lower than in existing applications, the EoR signal and foreground emission as well as the beam effects are stationary or semi-stationary.

In this paper, a novel deep-learning-based method that uses a CDAE is proposed to tackle the complicated frequency-dependent beam effects and to separate the EoR signal along the frequency dimension. In [Section 2](#), we introduce the CDAE and elaborate the proposed method. In [Section 3](#), we demonstrate the performance of the CDAE by applying it to the simulated SKA images. We discuss the method and carry out [comparisons to traditional methods](#) in [Section 4](#). Finally, we summarise our work in [Section 5](#). The implementation code and data are made public at <https://github.com/lwieitianux/cdae-eor>.

## 2 METHODOLOGY

### 2.1 Convolutional denoising autoencoder

Autoencoders are a type of neural networks that aim at learning useful features from the input data in an unsupervised manner and are usually used for dimensionality reduction (e.g., [Hinton & Salakhutdinov 2006](#); [Wang et al. 2014](#)) and data denoising (e.g., [Xie et al. 2012](#); [Lu et al. 2013](#); [Bengio et al. 2013](#)). An autoencoder is composed of an encoder and a decoder, which can be characterised by the functions  $f(\cdot)$  and  $g(\cdot)$ , respectively. The encoder maps the input  $\mathbf{x}$  to an internal code  $\mathbf{h}$ , i.e.,  $\mathbf{h} = f(\mathbf{x})$ , and the decoder tries to reconstruct the wanted signal from the code  $\mathbf{h}$ , i.e.,  $\mathbf{r} = g(\mathbf{h})$ , where  $\mathbf{x}$ ,  $\mathbf{h}$ , and  $\mathbf{r}$  are all vectors for this work. By placing constraints (e.g., dimensionality and sparsity) on the internal code  $\mathbf{h}$  and training the autoencoder to minimize the loss  $L(\mathbf{r}, \mathbf{x})$ , which quantifies the difference between the reconstruction  $\mathbf{r}$  and the input  $\mathbf{x}$ , the autoencoder is trained to learn the codes that effectively represent the input data (e.g., [Goodfellow et al. 2016](#), chapter 14).

In order to make the autoencoder learn a better representation of the input to achieve better performance, [Vincent et al. \(2008, 2010\)](#) proposed a novel training strategy based on the denoising

criterion, which states that ‘a good representation is the one that can be obtained robustly from a corrupted input and that will be useful for recovering the corresponding clean input.’ With this training strategy, we artificially corrupt the original input  $\mathbf{x}$  (e.g., by adding noise), feed the corrupted input  $\mathbf{x}'$  to the autoencoder, and then train it to reconstruct the original uncorrupted input by minimizing the loss  $L(\mathbf{r}, \mathbf{x})$ . During this denoising process, the autoencoder is required to extract features that capture useful structures in the original input, i.e., a good representation of the input. Since the trained autoencoder is able to denoise the corrupted input  $\mathbf{x}'$  to reconstruct the original uncorrupted input  $\mathbf{x}$ , it is hence called a ‘denoising autoencoder’ and is commonly applied to denoise images and other signals (e.g., [Xie et al. 2012](#); [Bengio et al. 2013](#); [Gondara 2016](#)).

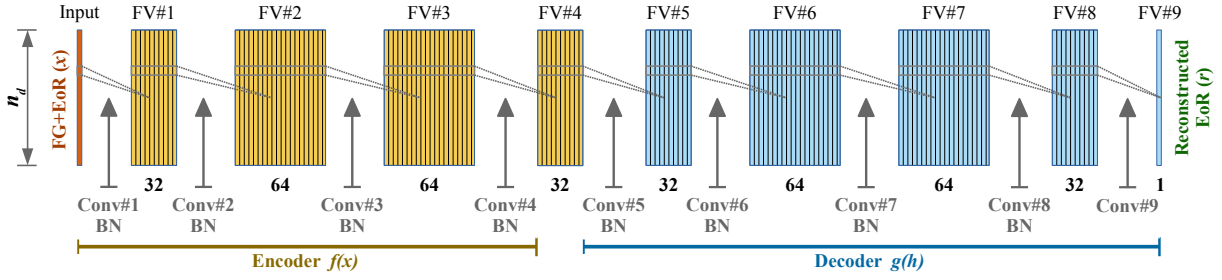
Classic autoencoders are built with fully connected layers, each neuron of which is connected with every neuron in the previous layer. This makes the total number of parameters increase exponentially with the number of layers and also forces the extracted features to be global, which is suboptimal to represent the input (e.g., [Masci et al. 2011](#)). On the other hand, convolutional layers as used in convolutional neural networks (CNNs) can well exploit the local features in the data by making use of a set of small filters and sharing their weights among all locations in the data (e.g., [Lecun et al. 1998](#)). Therefore, CNNs generally have two or more orders of magnitude less parameters than the analogous fully connected neural networks (e.g., [Grais & Plumbley 2017](#)), which greatly reduces the resource usages (e.g., memory, time) for training. Many convolutional layers can also be easily stacked to build highly expressive CNNs that achieve outstanding performance in image classification and relevant fields (e.g., [Krizhevsky et al. 2012](#); [Simonyan & Zisserman 2014](#); [Szegedy et al. 2015](#); [Ma et al. 2018](#)). By utilising multiple convolutional layers instead of fully connected layers in a denoising autoencoder, the CDAE gains the powerful feature extraction capabilities of CNNs, which helps improve its denoising performance, and can reconstruct even seriously corrupted signal (e.g., [Du et al. 2017](#)). In consequence, the CDAE may be well suited to exploit the complicated differences between the EoR signal and the foreground emission to reconstruct the EoR signal.

### 2.2 Network architecture

The CDAE proposed in this work consists of multiple convolutional layers that are divided into the encoder and decoder parts, between which there is no strict boundary because we focus on the feature extraction and denoising capabilities of the CDAE rather than on the specific formats of the internal codes  $\mathbf{h}$ . To some degree, the proposed CDAE is analogous to a CNN without the fully connected layers needed for classification. For the  $l$ -th convolutional layer that has  $m_l$  filters, a set of  $m_l$  feature vectors  $\left\{ \mathbf{v}_i^{(l)} \right\}; i = 1, 2, \dots, m_l$  are generated as the output of this layer by convolving the output of the previous layer  $\left\{ \mathbf{v}_j^{(l-1)} \right\}; j = 1, 2, \dots, m_{l-1}$  with each of the filters, i.e.,

$$\mathbf{v}_i^{(l)} = \phi^{(l)} \left( \sum_{j=1}^{m_{l-1}} \mathbf{v}_j^{(l-1)} * \mathbf{W}_i^{(l)} + b_i^{(l)} \right), \quad i = 1, 2, \dots, m_l, \quad (1)$$

where  $\mathbf{W}_i^{(l)}$  and  $b_i^{(l)}$  are the weights and bias of the  $i$ -th filter in the  $l$ -th layer,  $\phi^{(l)}(\cdot)$  is the layer’s activation function, and ‘\*’ denotes the convolution operation. In this way, sophisticated higher-level features can be obtained by composing the lower-level ones extracted



**Figure 1.** The architecture of the proposed CDAE that consists of a 4-layer encoder and a 5-layer decoder. The orange and blue boxes represent the feature vectors (FV) generated by the encoder and decoder layers, respectively. The numbers marked below the boxes are the number of filters in the corresponding convolutional layers. The batch normalisation (BN) technique is applied to all layers except for the last layer.

by the previous layers, implementing a powerful feature extraction technique (LeCun et al. 2015).

Following the common practices (e.g., Géron 2017; Suganuma et al. 2018), we adopt filters of size 3 in all layers and use the exponential linear unit (ELU; Clevert et al. 2016) as the activation function  $\phi^{(l)}(\cdot)$  for all layers except for the last layer, which uses the hyperbolic tangent function (i.e., tanh; see also Section 3.2). In addition, the batch normalisation technique is applied to all layers except for the last layer to improve the training process as well as to act as a regularizer to prevent over-fitting (Ioffe & Szegedy 2015).

To determine the number of convolutional layers and the number of filters in each layer, we have tested multiple CDAE architectures, each containing a different number of layers and filters. After evaluating their performances (see also Section 3.3), the simplest one with sufficiently good performance is selected, which consists of a 4-layer encoder with (32, 64, 64, 32) filters and a 5-layer decoder with (32, 64, 64, 32, 1) filters, as illustrated in Fig. 1. We note that the pooling and upsampling layers are not used in the CDAE because they have very little impact on the performance according to our tests (see also Springenberg et al. 2015).

### 2.3 Training and evaluation

At the beginning, the parameters (i.e., the weights and biases of filters in all layers) of the CDAE are initialised randomly according to the He uniform initialiser (He et al. 2015). To obtain an effective CDAE, these parameters need to be trained using the training data. The following three datasets are required to train and evaluate the CDAE (e.g., Ripley 1996): (1) training set ( $S_{\text{tr}}$ ) to fit the parameters by minimizing the loss  $L$ ; (2) validation set ( $S_{\text{val}}$ ) to validate the training process as well as to help determine the hyperparameters (e.g., the number of layers and filters; Section 2.2); (3) test set ( $S_{\text{test}}$ ) that is solely used to evaluate the performance of the trained CDAE. Each dataset is collection of many data points of  $(\mathbf{x}, \mathbf{x}_{\text{eor}})$ , where  $\mathbf{x} = \mathbf{x}_{\text{eor}} + \mathbf{x}_{\text{fg}}$  is the total emission of one sky pixel to be fed into the CDAE, and  $\mathbf{x}_{\text{eor}}$  is the corresponding EoR signal in the same pixel.

During each training epoch, the total emission  $\mathbf{x}^{(i)}$  is fed into the CDAE and goes through all the convolutional layers (Eq. 1) to derive the CDAE output  $\mathbf{r}_{\text{eor}}^{(i)}$ , which is expected to be the reconstructed EoR signal. The difference between the reconstructed EoR signal  $\mathbf{r}_{\text{eor}}^{(i)}$  and the input EoR signal  $\mathbf{x}_{\text{eor}}^{(i)}$  is the loss  $L$  and can be quantified using the mean squared error (MSE), i.e.,

$$L = \frac{1}{N_{\text{tr}}} \sum_{i=1}^{N_{\text{tr}}} \left[ \mathbf{r}_{\text{eor}}^{(i)} - \mathbf{x}_{\text{eor}}^{(i)} \right]^T \left[ \mathbf{r}_{\text{eor}}^{(i)} - \mathbf{x}_{\text{eor}}^{(i)} \right], \quad (2)$$

where  $N_{\text{tr}}$  is the number of data points in the training set  $S_{\text{tr}}$ . By

applying the back-propagation method (e.g., Rumelhart et al. 1986; LeCun et al. 1998), the loss  $L$  is then propagated backward in the CDAE to calculate the gradients  $(\partial L / \partial \theta)$  with respect to every parameter  $\theta$ , which are used to update the parameters to reduce the loss  $L$ , i.e., to improve the quality of the reconstructed EoR signal. As the training goes for more epochs, the initially randomised CDAE is gradually shaped into a tailored network targeting at the EoR signal separation task. In other words, the CDAE, driven by the training data, is learning to better extract the features of the EoR signal, from which the EoR signal can be reconstructed more accurately. Along the training process, the CDAE is also evaluated on the validation set ( $S_{\text{val}}$ ) to validate whether the training process goes well (e.g., no over-fitting). However, the validation set and its evaluation results should not be used to update the parameters (e.g., Russell & Norvig 2009).

To evaluate the performance of the CDAE in reconstructing the EoR signal, the commonly used Pearson's correlation coefficient (e.g., Harker et al. 2009; Chapman et al. 2013) is adopted to measure the similarity between the reconstructed EoR signal  $\mathbf{r}_{\text{eor}}$  and the input EoR signal  $\mathbf{x}_{\text{eor}}$ :

$$\rho(\mathbf{r}_{\text{eor}}, \mathbf{x}_{\text{eor}}) = \frac{\sum_{j=1}^n (r_{\text{eor},j} - \bar{r}_{\text{eor}})(x_{\text{eor},j} - \bar{x}_{\text{eor}})}{\sqrt{\sum_{j=1}^n (r_{\text{eor},j} - \bar{r}_{\text{eor}})^2 \sum_{j=1}^n (x_{\text{eor},j} - \bar{x}_{\text{eor}})^2}}, \quad (3)$$

where  $n$  is the length of the signals, and  $\bar{r}_{\text{eor}}$  and  $\bar{x}_{\text{eor}}$  represent the mean values. The closer to one the correlation coefficient is, the better the separation performance.

## 3 EXPERIMENTS

### 3.1 Simulation of the SKA images

We carry out end-to-end simulations to generate the SKA images to train the proposed CDAE and evaluate its performance. A representative frequency band, namely 154–162 MHz, is chosen as an example (e.g., Datta et al. 2010) and is divided into  $n_f = 101$  channels with a resolution of 80 kHz. At each frequency channel, the sky maps of the foreground emission and the EoR signal are simulated within an area of  $10^\circ \times 10^\circ$  and are pixelized into  $1800 \times 1800$  with a pixel size of  $20''$ .

Based on our previous work (Wang et al. 2010), we simulate the foreground emission by taking into account the contributions from the Galactic synchrotron and free-free radiations, extragalactic point sources, and radio haloes. The Galactic synchrotron radiation is simulated by extrapolating the Haslam 408 MHz map with a power-law spectrum, the index of which is given by the synchrotron spectral index map (Giardino et al. 2002) to account for its variation with sky

positions. The reprocessed Haslam 408 MHz map<sup>1</sup> (Remazeilles et al. 2015), which has significantly better instrument calibration and more accurate extragalactic source subtraction, is used as the template to obtain enhanced simulation results over Wang et al. (2010). By employing the tight relation between the H $\alpha$  and free-free emissions (see Dickinson et al. 2003, and references therein), the Galactic free-free radiation can be derived from the H $\alpha$  survey map (Finkbeiner 2003). Since the Galactic diffuse emissions vary remarkably across the sky, we simulate them at a central position of (R.A., Dec.) = (0°, -27°), which has a high galactic latitude ( $b = -78.5^\circ$ ) and is an appropriate choice for the simulation of SKA images (e.g., Beardsley et al. 2016). We account for the following 5 types of extragalactic point sources: (1) star-forming and starburst galaxies, (2) radio-quiet active galactic nuclei (AGNs), (3) Fanaroff-Riley type I and type II AGNs, (4) GHz-peaked spectrum AGNs, and (5) compact steep spectrum AGNs. The former three types of sources are simulated by utilising the data published by Wilman et al. (2008) and the latter two types are simulated by employing their corresponding luminosity functions and spectral models. Similar to the real-time peeling of the brightest point sources in practical data analysis pipelines (e.g., Mitchell et al. 2008; Intema et al. 2009), we assume that sources with a 158 MHz flux density  $S_{158} > 10$  mJy have been removed (e.g., Liu et al. 2009a). The radio haloes are simulated by generating a sample of galaxy clusters with the Press-Schechter formalism (Press & Schechter 1974) and then applying multiple scaling relations (e.g., between cluster mass and X-ray temperature, between X-ray temperature and radio power) to derive their radio emissions.

In regard to the simulation of the EoR signal, we take advantage of the 2016 data release from the Evolution Of 21 cm Structure project<sup>2</sup> (Mesinger et al. 2016) and extract the image slices at corresponding redshifts (i.e., frequencies) from the light-cone cube of the recommended ‘faint galaxies’ case. The extracted image slices are then re-scaled to match the sky coverage and pixel size of the foreground maps.

To incorporate the realistic frequency-dependent beam effects into the simulated sky maps, we further adopt the latest SKA1-Low layout configuration<sup>3</sup> to simulate instrument observations. The SKA1-Low interferometer is composed of 512 stations, each of which contains 256 antennas randomly distributed inside a circle of 35 m in diameter. The 512 stations are divided into two parts: (1) 224 stations are randomly distributed within the ‘core’ region of 1000 m in diameter; (2) the remaining stations are placed on 3 spiral arms extending up to a radius of about 35 km. For each sky map, we employ the OSKAR<sup>4</sup> simulator (Mort et al. 2010) to perform 6-hour synthesis imaging, yielding the visibility data, from which the ‘observed’ image is created by the WSCLEAN<sup>5</sup> imager (Offringa et al. 2014). In order to emphasize the faint and relatively diffuse EoR signal, the natural weighting and baselines of 30–1000 wavelengths are utilised in the imaging process. Finally, the created images are cropped to keep only the central  $2^\circ \times 2^\circ$  regions (i.e.,  $360 \times 360$  pixels) for the purpose of the best quality. Therefore, we obtain a

pair of image cubes of size  $360 \times 360 \times 101$  for the EoR signal ( $C_{\text{eor}}^{(1)}$ ) and the foreground emission ( $C_{\text{fg}}^{(1)}$ ), respectively (see Fig. 2 for the simulated images at the central frequency of 158 MHz). To better illustrate the impacts of beam effects on the foreground spectra, we show in Fig. 3 the root-mean-square (r.m.s.) spectra and the corresponding differential spectra of the foreground image cubes with and without the beam effects considered. Compared to the ideal sky foreground, the spectral smoothness of the ‘observed’ foreground is seriously damaged by the rapid fluctuations resulted from the beam effects. Such fluctuations have scales of  $< 1$  MHz and look more similar to the EoR signal, which can impose a great difficulty on traditional foreground removal methods (e.g., Liu et al. 2009a).

Considering that the training and evaluation of the CDAE require three datasets (i.e., training, validation, and test; Section 2.3), if there are only one pair of image cubes, the test set could only contain a small fraction of all the pixels that are randomly distributed on the sky, from which it is impossible to obtain a complete image of the reconstructed EoR signal. Consequently, it is beneficial to simulate the second pair of image cubes that are solely used as the test set. To this end, we simulate the Galactic diffuse radiations at a central coordinate of (R.A., Dec.) = ( $3^\circ$ ,  $-27^\circ$ ), i.e.,  $3^\circ$  away from the first pair of image cubes, which is sufficient because the finally cropped image cubes only cover a sky area of  $2^\circ \times 2^\circ$ . Since extragalactic point sources, radio haloes, and the EoR signal are mostly isotropic, we shift their sky maps simulated above by  $3^\circ$  to generate the new sky maps. Following the same procedures to simulate instrument observations, we obtain the second pair of image cubes ( $C_{\text{eor}}^{(2)}, C_{\text{fg}}^{(2)}$ ).

We note that the simulations do not include thermal noise because the proposed method is designed to create tomographic EoR images from very deep SKA observations that have a sufficiently low noise level. Among the 13 high priority science objectives identified for SKA phase I, the EoR imaging objective is of the highest priority<sup>6</sup>. It is planned to select 5 clean fields on the sky and to observe each of them for about 1000 h, reaching an unprecedented image noise level of  $\lesssim 1$  mK that allows to directly image the reionization structures (e.g., Mellema et al. 2013, 2015; Koopmans et al. 2015).

### 3.2 Data preprocessing

The dataset  $S = \{(\mathbf{x}, \mathbf{x}_{\text{eor}})\}$  for the CDAE is derived from the simulated image cubes  $C_{\text{eor}}$  and  $C_{\text{fg}}$ , each data point ( $\mathbf{x} = \mathbf{x}_{\text{eor}} + \mathbf{x}_{\text{fg}}$ ,  $\mathbf{x}_{\text{eor}}$ ) representing the total emission and the EoR signal of one sky pixel, respectively. The dataset thus has  $N_S = 360 \times 360 \times 2 = 259\,200$  data points in total.

For the input data  $X = \{\mathbf{x}\}$ , we propose to apply the Fourier Transform (FT) along the frequency dimension, which makes the EoR signal more distinguishable from the foreground emission and thus easier to be learned by the CDAE (a comparison with the results derived without applying the FT is presented in Section 4.1). The Blackman-Nuttall window function is applied to suppress the FT side-lobes caused by the sharp discontinuities at both ends of the finite frequency band (e.g., Chapman et al. 2016). It is sufficient to keep only half the Fourier coefficients because  $\mathbf{x}$  is real, thus  $\mathbf{x}$  of length  $n_f = 101$  is transformed to be 51 complex Fourier

<sup>1</sup> The reprocessed Haslam 408 MHz map: [http://www.jb.man.ac.uk/research/cosmos/haslam\\_map/](http://www.jb.man.ac.uk/research/cosmos/haslam_map/)

<sup>2</sup> Evolution Of 21 cm Structure: <http://homepage.sns.it/mesinger/EOS.html>


<sup>3</sup> SKA1-Low layout: [https://astronomers.skatelescope.org/wp-content/uploads/2016/09/SKA-TEL-SK0-0000422\\_02\\_SKA1\\_LowConfigurationCoordinates-1.pdf](https://astronomers.skatelescope.org/wp-content/uploads/2016/09/SKA-TEL-SK0-0000422_02_SKA1_LowConfigurationCoordinates-1.pdf)

<sup>4</sup> OSKAR: <https://github.com/OxfordSKA/OSKAR> (version 2.7.0)

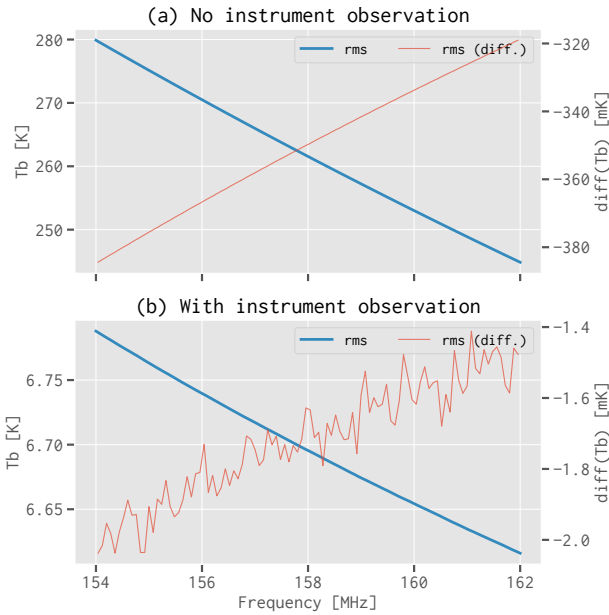
<sup>5</sup> WSClean: <https://sourceforge.net/p/wsclean> (version 2.5)

<sup>6</sup> SKA1 level 0 science requirements: [https://www.skatelescope.org/wp-content/uploads/2014/03/SKA-TEL-SK0-0000007\\_SKA1\\_Level\\_0\\_Science\\_RequirementsRev02-part-1-signed.pdf](https://www.skatelescope.org/wp-content/uploads/2014/03/SKA-TEL-SK0-0000007_SKA1_Level_0_Science_RequirementsRev02-part-1-signed.pdf)





**Figure 2.** Simulated images of the EoR signal (left panel) and the foreground emission (right panel) at 158 MHz. Both images have a size of  $360 \times 360$  and cover a sky area of  $2^\circ \times 2^\circ$ . The blobs in the right panel show the bright point sources and radio haloes.



**Figure 3.** The r.m.s. spectra (the bold blue lines) and the corresponding differential spectra (the thin red lines) of the foreground image cubes. The top and bottom panels show the cases without and with instrument observation, respectively.

coefficients. The  $n_{\text{ex}}$  coefficients of the lowest Fourier frequencies are excised since they are mostly contributed by the spectral-smooth foreground emission. We adopt  $n_{\text{ex}} = 6$  to achieve a balance between the foreground suppression and the loss of the EoR signal. The real and imaginary parts of the remaining 45 complex coefficients are then concatenated into a new real vector of length  $n_d = 90$ , since the CDAE requires real data. Finally, the data are zero-centred and normalised to have unit variance.

The preprocessing steps for the input EoR signal  $X_{\text{eor}} = \{\mathbf{x}_{\text{eor}}\}$  are basically the same except for minor adjustments. After applying the FT, excising the  $n_{\text{ex}}$  lowest Fourier components, and concatenating the real and imaginary parts, the data elements that have a

value less than the 1<sup>st</sup> percentile or greater than the 99<sup>th</sup> percentile are truncated, in order to prevent the possible outliers hindering the training of the CDAE. Finally, the value range of the data is scaled to be  $[-1, 1]$  by dividing by the maximum absolute value, which allows to use the ‘tanh’ activation function whose value range is also  $[-1, 1]$  in the output layer of the CDAE (Section 2.2).

### 3.3 Training and results

The preprocessed data of the first cube pair  $(C_{\text{eor}}^{(1)}, C_{\text{fg}}^{(1)})$  are randomly partitioned into the training set ( $S_{\text{tr}}$ ; corresponding to 80 per cent of the pixels, or 103 680 data points) and the validation set ( $S_{\text{val}}$ ; 20 per cent, or 25 920 data points). The preprocessed data of the second cube pair  $(C_{\text{eor}}^{(2)}, C_{\text{fg}}^{(2)})$  are solely used as the test set ( $S_{\text{test}}$ ; 129 600 data points). The use of the whole image cubes as the test set enables us to create complete images of the reconstructed EoR signal.


We implement the proposed CDAE using the KERAS<sup>7</sup> framework (Chollet et al. 2015) with the TENSORFLOW<sup>8</sup> back end (Abadi et al. 2015), which is accelerated by the CUDA<sup>9</sup> toolkit. We adopt a small initial learning rate ( $\alpha = 10^{-5}$ ) and use the Adam optimisation method (Kingma & Ba 2015). The CDAE is trained on the training set ( $S_{\text{tr}}$ ) with a batch size of 100 until the training loss converges, which takes about 50 epochs.

The training and validation losses together with the evaluation index (i.e., the correlation coefficient  $\rho$ ) calculated on the validation set  $S_{\text{val}}$  during the training phase are shown in Fig. 4. The steadily decreasing losses and increasing correlation coefficient suggest that the CDAE is well trained without over-fitting. After training, the evaluation with the test set  $S_{\text{test}}$  yields a high correlation coefficient of  $\bar{\rho}_{\text{CDAE}} = 0.929 \pm 0.045$  between the reconstructed and input EoR signals, which indicates that the trained CDAE achieves excellent performance in reconstructing the EoR signal. As an example,

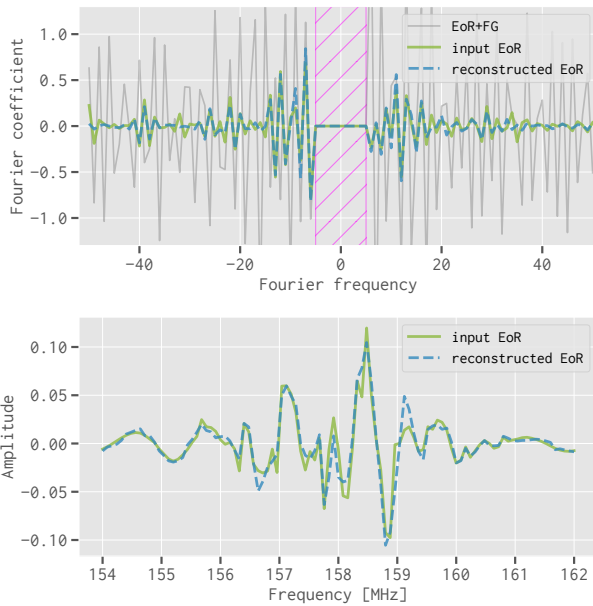
<sup>7</sup> Keras: <https://keras.io> (version 2.2.4)

<sup>8</sup> TensorFlow: <https://www.tensorflow.org> (version 1.12.0)

<sup>9</sup> CUDA: <https://developer.nvidia.com/cuda-zone> (version 9.1.85)



**Figure 4.** The training loss (the solid red line), validation loss (the solid blue line), and correlation coefficient ( $\rho$ ) calculated on the validation set  $S_{\text{val}}$  (the dashed blue line with the shaded region representing the standard deviation) along the training of the CDAE.



**Figure 5.** An example of the EoR signal reconstructed by the trained CDAE for one pixel in  $S_{\text{test}}$ . **(top)** The input EoR signal  $\mathbf{x}_{\text{eor}}$  (the solid green line) and the reconstructed EoR signal  $\mathbf{r}_{\text{eor}}$  (the dashed blue line) in the Fourier domain. The correlation coefficient between the input and reconstructed EoR signals is  $\rho = 0.931$ . The grey line represents the input total emission  $\mathbf{x} = \mathbf{x}_{\text{fg}} + \mathbf{x}_{\text{eor}}$ . The magenta hatched region marks the excised Fourier coefficients in data preprocessing. **(bottom)** The input EoR signal  $\mathbf{x}_{\text{eor}}$  (the solid green line) and the reconstructed EoR signal  $\mathbf{r}_{\text{eor}}$  (the dashed blue line) transformed back to the observing frequency domain.

Fig. 5 illustrates the reconstructed EoR signal ( $\rho = 0.931$ ) for one pixel in  $S_{\text{test}}$ .

Since the test set  $S_{\text{test}}$  is derived from the whole image cubes  $(C_{\text{eor}}^{(2)}, C_{\text{fg}}^{(2)})$ , we are able to create complete images of the reconstructed EoR signal and calculate the corresponding power spectrum. Taking the input and reconstructed EoR images at the central frequency of 158 MHz as an example (Fig. 6), the reconstructed EoR

signal exhibits almost identical structures and amplitudes as the input EoR signal. We note that the reconstructed EoR image has weak but detectable redundant ripples on scales of about 10 pixels (i.e.,  $\sim 200''$ ), which are associated with the excision of the  $n_{\text{ex}}$  lowest Fourier frequencies in data preprocessing (Section 3.2). In addition, we calculate the two-dimensional power spectra from the image cubes of the input and reconstructed EoR signals (Fig. 7). It illustrates that the trained CDAE well recovers the EoR signal on all covered scales except for a very thin stripe region at  $k_{\perp} \approx 0.7 \text{ Mpc}^{-1}$ , where extra powers are generated by the aforementioned ripples in the reconstructed EoR images. We also note that there is a barely visible line at  $k_{\perp} \approx 0.1 \text{ Mpc}^{-1}$  in both power spectra, which is caused by the boundary effect of Fourier transforming the finite frequency band.

The results clearly demonstrate that the trained CDAE is able to well reconstruct the EoR signal, overcoming the complicated beam effects. The achieved excellent performance of the CDAE can be mainly attributed to the architecture of stacking multiple convolutional layers, which implements a powerful feature extraction technique by hierarchically combining the basic features learned in each layer to build more and more sophisticated features (LeCun et al. 2015). Combined with the flexibility provided by the 53 569 trainable parameters, the CDAE, after being well trained, can intelligently learn a model that is optimised to accurately separate the faint EoR signal (e.g., Domingos 2012).

### 3.4 Further validation of the CDAE

With the purpose of further validating that the trained CDAE has actually learned the useful features of the EoR signal, we employ the occlusion method (Zeiler & Fergus 2014) to visualise the sensitivity of the trained CDAE to the different part of the input data. At each time, we occlude 3 adjacent elements of every input total emission  $\mathbf{x}$  in the validation set  $S_{\text{val}}$ , and then measure the CDAE's sensitivity to the occluded part, which is calculated as the occlusion-induced performance loss, i.e.,


$$s = \frac{1}{N_{\text{val}}} \sum_{i=1}^{N_{\text{val}}} \left[ \rho(\mathbf{r}_{\text{eor}}^{(i)}, \mathbf{x}_{\text{eor}}^{(i)}) - \rho(\mathbf{R}_{\text{eor}}^{(i)}, \mathbf{x}_{\text{eor}}^{(i)}) \right], \quad (4)$$

where  $N_{\text{val}}$  is the number of data points in the validation set,  $\mathbf{x}_{\text{eor}}^{(i)}$  is the input EoR signal, and  $\mathbf{r}_{\text{eor}}^{(i)}$  and  $\mathbf{R}_{\text{eor}}^{(i)}$  are the reconstructed EoR signals without and with applying the occlusion, respectively. By varying the occlusion part of the input data and calculating the sensitivities, we obtain the CDAE's sensitivity distribution ( $s$ ) to every part of the input data, as shown in Fig. 8, where the r.m.s. amplitudes of the foreground emission ( $\mathbf{y}_{\text{fg}}$ ) and the EoR signal ( $\mathbf{y}_{\text{eor}}$ ) are also plotted. We find that the sensitivity distribution is significantly more correlated with the EoR signal [ $\rho(s, \mathbf{y}_{\text{eor}}) = 0.742$ ] than the foreground [ $\rho(s, \mathbf{y}_{\text{fg}}) = 0.562$ ]. This verifies that the trained CDAE has learned useful features of the EoR signal to distinguish it from the foreground emission and thus becomes more sensitive to the data parts of higher signal-to-noise ratio.


## 4 DISCUSSIONS

### 4.1 Why preprocess the dataset with Fourier Transform?

We perform another experiment using the same CDAE architecture, datasets, and data preprocessing steps, except for applying the FT



**Figure 6.** Comparison between the input EoR image (the left panel) and reconstructed EoR image (the right panel) at the central frequency of 158 MHz. The images have the same size ( $360 \times 360$  pixel) and the figures share the same colour bar (the amplitude is normalised for the CDAE).




**Figure 7.** Comparison of two-dimensional power spectra between the input (the left panel) and reconstructed (the right panel) EoR signals.

as depicted in Section 3.2. After training the CDAE in the same way as described in Section 3.3, the correlation coefficient between the reconstructed and input EoR signals evaluated on the test set  $S_{\text{test}}$  reaches only  $\bar{\rho}_{\text{noft}} = 0.628 \pm 0.167$ , which indicates a significantly worse performance compared to the case with FT applied. As presented in Fig. 9, the training loss decreases more slowly and converges after about 100 epochs. We also find that the training process is slightly unstable given the small spikes on the curves of both the loss and correlation coefficient. These indicate that it is beneficial to preprocess the dataset by applying the FT along the frequency dimension, because the EoR signal and the foreground emission become more distinguishable in the Fourier domain, where the fluctuating EoR signal concentrates on larger Fourier modes while the spectral-smooth foreground emission distributes mainly on smaller Fourier modes (e.g., Parsons et al. 2012).


#### 4.2 Comparing to traditional methods

A variety of methods have been proposed to remove the foreground contamination with the aim of revealing the faint EoR signal. These methods can be broadly classified into two categories: (1) parametric methods that apply a parametric model (e.g., a low-degree polynomial) to fit and remove the foreground emission (e.g., Wang et al. 2006; Jelić et al. 2008; Liu et al. 2009b; Wang et al. 2013; Bonaldi & Brown 2015); (2) non-parametric methods, which do not assume a specific parametric model for the foreground emission but exploit the differences between the foreground emission and the EoR signal (e.g., their different spectral features) to separate them (e.g., Harker et al. 2009; Chapman et al. 2012, 2013; Gu et al. 2013; Mertens et al. 2018).

In order to further demonstrate the performance of our method, we compare it to two representative traditional methods: the polynomial fitting method (e.g., Wang et al. 2006) and the continuous wavelet transform (CWT) method (Gu et al. 2013). The polynomial



**Figure 8.** The CDAE’s sensitivity distribution  $s$  (the blue lines in both panels) obtained by applying the occlusion method. We also plot the r.m.s. amplitudes of the foreground emission ( $y_{fg}$ , the red line in the left panel) and the EoR signal ( $y_{eor}$ , the green line in the right panel). The sensitivity distribution  $s$  is more correlated with the EoR signal [ $\rho(s, y_{eor}) = 0.742$ ] than the foreground [ $\rho(s, y_{fg}) = 0.562$ ].



**Figure 9.** Same as Fig. 4 but for the case that the data are preprocessed without applying the FT.

fitting method is the best representative of the parametric methods because it is widely used due to its simplicity and robustness (e.g., Jelić et al. 2008; Liu et al. 2009a; Pritchard & Loeb 2010) and has been compared to various other foreground removal methods (e.g., Harker et al. 2009; Alonso et al. 2015; Chapman et al. 2015). Among the non-parametric category, the CWT method is chosen since it performs similarly well as other non-parametric methods, such as the Wp smoothing method (Harker et al. 2009) and the generalized morphological component analysis method (Chapman et al. 2013), meanwhile it is faster and simpler (Gu et al. 2013; Chapman et al. 2015).

With the polynomial fitting method, a low-degree polynomial is fitted along the frequency dimension for each sky pixel in the image cube of the total emission (i.e.,  $C_{tot} = C_{eor} + C_{fg}$ ). Then by subtracting the fitted smooth component, which is regarded as the foreground emission, the EoR signal is expected to be uncovered. Using the same image cubes ( $C_{eor}^{(2)}, C_{fg}^{(2)}$ ) simulated in Section 3.1, we have tested polynomials of the degree from 2 (quadratic) to 5 (quintic), and find that the quartic polynomial (degree of 4) can give the best result. However, the correlation coefficient calculated for the

separated EoR signal in such a case is only  $\bar{\rho}_{poly} = 0.296 \pm 0.121$ , which indicates that the polynomial fitting method performs poorly in removing the foreground emission.

The CWT method works based on the same assumption as other foreground removal methods that the foreground emission is spectrally smooth while the EoR signal fluctuates rapidly along the frequency dimension. After applying the CWT, the foreground emission and the EoR signal locate at different spectral scales. Therefore, the foreground emission can be easily separated from the EoR signal and be removed (Gu et al. 2013). For each sky pixel, the spectrum of the total emission is transformed into the wavelet space by applying the CWT with the Morlet wavelet function. In the wavelet space, after identifying and removing the coefficients that are mainly contributed by the foreground emission, the remaining coefficients are transformed back to the frequency space to obtain the spectrum with the foreground emission removed, which is expected to be the EoR signal. By evaluating on the same dataset ( $C_{eor}^{(2)}, C_{fg}^{(2)}$ ), we have tuned the method parameters (minimum scale  $s_{min}$ , maximum scale  $s_{max}$ , number of scales  $n_s$ , and cone of influence  $c_i$ ) and adopt  $s_{min} = 7.4$ ,  $s_{max} = 50.0$ ,  $n_s = 50$ , and  $c_i = 1.6$  to obtain the relatively best performance, which is, however, only  $\bar{\rho}_{cwt} = 0.198 \pm 0.160$ . We note that the CWT method performs slightly worse than the polynomial fitting method, which is difference from the comparison in Gu et al. (2013). This may be caused by the more serious boundary effect since our simulated data have a narrower bandwidth and coarser frequency resolution than those of Gu et al. (2013).

The main reason that both traditional foreground removal methods only obtain remarkably inferior results is that the smoothness of the foreground spectra is seriously damaged by the frequency-dependent beam effects, which cause rapid fluctuations of strength the same order as the EoR signal on the originally smooth foreground spectra [Fig. 3(b)]. As a result, the foreground spectra complicated by the beam effects cannot be well fitted by a low-degree polynomial and have more similar spectral scales as the EoR signal. In consequence, both methods are unable to well model the complicated foreground spectra and thus have great difficulties in removing them. On the contrary, given its data-driven nature and powerful feature extraction capabilities, the CDAE is able to distil knowledge from the training data and learns the features to



distinguish the EoR signal from the fluctuations arising from the beam effects. Hence, the CDAE achieves superior performance in separating the EoR signal.

## 5 SUMMARY

The frequency-dependent beam effects of interferometers can cause rapid fluctuations along the frequency dimension, which destroy the smoothness of the foreground spectra and prevent traditional foreground removal methods from uncovering the EoR signal. Given the difficulties in crafting practicable models to overcome the complicated beam effects, methods that can intelligently learn tailored models from the data seem more feasible and appealing. To this end, we have proposed a deep-learning-based method that uses a 9-layer CDAE to separate the EoR signal. The CDAE has been trained on the simulated SKA images and has achieved excellent performance. We conclude that the CDAE has outstanding ability to overcome the complicated beam effects and accurately separate the faint EoR signal, exhibiting the great potential of deep-learning-based methods to play an important role in the forthcoming EoR experiments.

## ACKNOWLEDGEMENTS

We thank the reviewer and editor for their useful comments that greatly help improve the manuscript. We also thank Jeffrey Hsu for reading the manuscript and providing helpful suggestions. This work is supported by the Ministry of Science and Technology of China (grant Nos. 2018YFA0404601, 2017YFF0210903), and the National Natural Science Foundation of China (grant Nos. 11433002, 11621303, 11835009, 61371147).

## REFERENCES

Abadi M., et al., 2015, TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems, <https://www.tensorflow.org/>  
 Alonso D., Bull P., Ferreira P. G., Santos M. G., 2015, *MNRAS*, **447**, 400  
 Beardsley A. P., et al., 2016, *ApJ*, **833**, 102  
 Bengio Y., Yao L., Alain G., Vincent P., 2013, in Proceedings of the 26th International Conference on Neural Information Processing Systems (NIPS 2013). Curran Associates Inc., USA, pp 899–907, <http://dl.acm.org/citation.cfm?id=2999611.2999712>  
 Bonaldi A., Brown M. L., 2015, *MNRAS*, **447**, 1973  
 Chapman E., et al., 2012, *MNRAS*, **423**, 2518  
 Chapman E., et al., 2013, *MNRAS*, **429**, 165  
 Chapman E., et al., 2015, Advancing Astrophysics with the Square Kilometre Array (AASKA14), p. 5  
 Chapman E., Zaroubi S., Abdalla F. B., Dulwich F., Jelić V., Mort B., 2016, *MNRAS*, **458**, 2928  
 Chollet F., et al., 2015, Keras, <https://keras.io>  
 Clevert D.-A., Unterthiner T., Hochreiter S., 2016, in The International Conference on Learning Representations (ICLR 2016). (arXiv:1511.07289)  
 Datta A., Bowman J. D., Carilli C. L., 2010, *ApJ*, **724**, 526  
 Dickinson C., Davies R. D., Davis R. J., 2003, *MNRAS*, **341**, 369  
 Domingos P., 2012, *Communications of the ACM*, **55**, 78  
 Du B., Xiong W., Wu J., Zhang L., Zhang L., Tao D., 2017, *IEEE Transactions on Cybernetics*, **47**, 1017  
 Finkbeiner D. P., 2003, *ApJS*, **146**, 407  
 Furlanetto S. R., 2016, *Understanding the Epoch of Cosmic Reionization: Challenges and Progress*, **423**, 247

Géron A., 2017, Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, 1st edn. O'Reilly Media, Inc.  
 Giardino G., Banday A. J., Górski K. M., Bennett K., Jonas J. L., Tauber J., 2002, *A&A*, **387**, 82  
 Gondara L., 2016, in 2016 IEEE 16th International Conference on Data Mining Workshops (ICDMW). pp 241–246 (arXiv:1608.04667), doi:10.1109/ICDMW.2016.0041  
 Goodfellow I., Bengio Y., Courville A., 2016, Deep Learning. MIT Press, <http://www.deeplearningbook.org>  
 Grais E. M., Plumbley M. D., 2017, in 5th IEEE Global Conference on Signal and Information Processing (GlobalSIP 2017). IEEE, pp 1265–1269 (arXiv:1703.08019), doi:10.1109/GlobalSIP.2017.8309164  
 Gu J., Xu H., Wang J., An T., Chen W., 2013, *ApJ*, **773**, 38  
 Harker G., et al., 2009, *MNRAS*, **397**, 1138  
 He K., Zhang X., Ren S., Sun J., 2015, in Proceedings of the 2015 IEEE International Conference on Computer Vision (ICCV 2015). IEEE Computer Society, Washington, DC, USA, pp 1026–1034, doi:10.1109/ICCV.2015.123  
 Herbel J., Kacprzak T., Amara A., Refregier A., Lucchi A., 2018, *Journal of Cosmology and Astro-Particle Physics*, **2018**, 054  
 Hinton G. E., Salakhutdinov R. R., 2006, *Science*, **313**, 504  
 Intema H. T., van der Tol S., Cotton W. D., Cohen A. S., van Bemmell I. M., Röttgering H. J. A., 2009, *A&A*, **501**, 1185  
 Ioffe S., Szegedy C., 2015, in Proceedings of the 32nd International Conference on International Conference on Machine Learning (ICML 2015). JMLR.org, pp 448–456  
 Jelić V., et al., 2008, *MNRAS*, **389**, 1319  
 Kingma D. P., Ba J., 2015, in The International Conference on Learning Representations (ICLR 2015). (arXiv:1412.6980)  
 Koopmans L., et al., 2015, Advancing Astrophysics with the Square Kilometre Array (AASKA14), p. 1  
 Krizhevsky A., Sutskever I., Hinton G. E., 2012, in Proceedings of the 25th International Conference on Neural Information Processing Systems (NIPS). Curran Associates Inc., USA, pp 1097–1105, <http://dl.acm.org/citation.cfm?id=2999134.2999257>  
 LeCun Y., Bottou L., Orr G. B., Müller K.-R., 1998, in Neural Networks: Tricks of the Trade. Springer-Verlag, London, UK, UK, pp 9–50, <http://dl.acm.org/citation.cfm?id=645754.668382>  
 LeCun Y., Bengio Y., Hinton G., 2015, *Nature*, **521**, 436  
 Lecun Y., Bottou L., Bengio Y., Haffner P., 1998, *Proceedings of the IEEE*, **86**, 2278  
 Liu A., Tegmark M., Zaldarriaga M., 2009a, *MNRAS*, **394**, 1575  
 Liu A., Tegmark M., Bowman J., Hewitt J., Zaldarriaga M., 2009b, *MNRAS*, **398**, 401  
 Lochner M., Natarajan I., Zwart J. T. L., Smirnov O., Bassett B. A., Oozeer N., Kunz M., 2015, *MNRAS*, **450**, 1308  
 Lu X., Tsao Y., Matsuda S., Hori C., 2013, in 14th Annual Conference of the International Speech Communication Association (INTERSPEECH 2013). pp 436–440, [https://www.isca-speech.org/archive/interspeech\\_2013/i13\\_0436.html](https://www.isca-speech.org/archive/interspeech_2013/i13_0436.html)  
 Ma Z., et al., 2018, arXiv e-prints,  
 Masci J., Meier U., Cireşan D., Schmidhuber J., 2011, in Proceedings of the 21th International Conference on Artificial Neural Networks (ICANN 2011). Springer-Verlag, pp 52–59, <http://dl.acm.org/citation.cfm?id=2029556.2029563>  
 Mellema G., et al., 2013, *Experimental Astronomy*, **36**, 235  
 Mellema G., Koopmans L., Shukla H., Datta K. K., Mesinger A., Majumdar S., 2015, Advancing Astrophysics with the Square Kilometre Array (AASKA14), p. 10  
 Mertens F. G., Ghosh A., Koopmans L. V. E., 2018, *MNRAS*, **478**, 3640  
 Mesinger A., Greig B., Sobacchi E., 2016, *MNRAS*, **459**, 2342  
 Mitchell D. A., Greenhill L. J., Wayth R. B., Sault R. J., Lonsdale C. J., Cappallo R. J., Morales M. F., Ord S. M., 2008, *IEEE Journal of Selected Topics in Signal Processing*, **2**, 707  
 Morales M. F., Wyithe J. S. B., 2010, *ARA&A*, **48**, 127  
 Mort B. J., Dulwich F., Salvini S., Adami K. Z., Jones M. E., 2010, in 2010 IEEE International Symposium on Phased Array Systems and

- Technology. pp 690–694, [doi:10.1109/ARRAY.2010.5613289](#)
- Offringa A. R., et al., 2014, *MNRAS*, **444**, 606
- Parsons A. R., Pober J. C., Aguirre J. E., Carilli C. L., Jacobs D. C., Moore D. F., 2012, *ApJ*, **756**, 165
- Press W. H., Schechter P., 1974, *ApJ*, **187**, 425
- Pritchard J. R., Loeb A., 2010, *Phys. Rev. D*, **82**, 023006
- Remazeilles M., Dickinson C., Banday A. J., Bigot-Sazy M.-A., Ghosh T., 2015, *MNRAS*, **451**, 4311
- Ripley B. D., 1996, Pattern Recognition and Neural Networks. Cambridge University Press, [doi:10.1017/CBO9780511812651](#)
- Rumelhart D. E., Hinton G. E., Williams R. J., 1986, *Nature*, **323**, 533
- Russell S., Norvig P., 2009, Artificial Intelligence: A Modern Approach, 3rd edn. Prentice Hall
- Shen H., George D., Huerta E. A., Zhao Z., 2017, preprint, ([arXiv:1711.09919](#))
- Simonyan K., Zisserman A., 2014, preprint, ([arXiv:1409.1556](#))
- Springenberg J. T., Dosovitskiy A., Brox T., Riedmiller M., 2015, in International Conference on Learning Representations (ICLR 2015). ([arXiv:1412.6806](#))
- Suganuma M., Ozay M., Okatani T., 2018, in Proceedings of the 35th International Conference on Machine Learning (ICML 2018). ([arXiv:1803.00370](#))
- Szegedy C., et al., 2015, in 2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2015). pp 1–9 ([arXiv:1409.4842](#)), [doi:10.1109/CVPR.2015.7298594](#)
- Vafaei Sadr A., Vos E. E., Bassett B. A., Hosenie Z., Oozeer N., Lochner M., 2018, preprint, ([arXiv:1807.02701](#))
- Vincent P., Larochelle H., Bengio Y., Manzagol P.-A., 2008, in Proceedings of the 25th International Conference on Machine Learning (ICML 2008). ACM, pp 1096–1103, [doi:10.1145/1390156.1390294](#)
- Vincent P., Larochelle H., Lajoie I., Bengio Y., Manzagol P.-A., 2010, The Journal of Machine Learning Research, **11**, 3371
- Wang X., Tegmark M., Santos M. G., Knox L., 2006, *ApJ*, **650**, 529
- Wang J., et al., 2010, *ApJ*, **723**, 620
- Wang J., et al., 2013, *ApJ*, **763**, 90
- Wang W., Huang Y., Wang Y., Wang L., 2014, in 2014 IEEE Conference on Computer Vision and Pattern Recognition Workshops. pp 496–503, [doi:10.1109/CVPRW.2014.79](#)
- Wilman R. J., et al., 2008, *MNRAS*, **388**, 1335
- Xie J., Xu L., Chen E., 2012, in Proceedings of the 25th International Conference on Neural Information Processing Systems (NIPS 2012). Curran Associates Inc., USA, pp 341–349
- Zeiler M. D., Fergus R., 2014, in Fleet D., Pajdla T., Schiele B., Tuytelaars T., eds, European Conference on Computer Vision (ECCV 2014). Springer, pp 818–833 ([arXiv:1311.2901](#)), [doi:10.1007/978-3-319-10590-1\\_53](#)

This paper has been typeset from a  $\text{\LaTeX}$  file prepared by the author.